


10 equal-count bins of time, 193–756 observations each



Mean over subjects at each time, in equal-count bins of time. Dark band: where the model puts 95% bin MEAN, across 200 generated datasets -- the reference the observed mean should sit inside. Light band: where it puts 95% of individual observations, which is the spread in the data rather than uncertainty about its mean. The mean band is much the narrower of the two, which is the point of the panel: a bias in the average level that disappears inside the observation cloud of ValueByTime is plain against it. LogLik is each row's log-likelihood, and its distribution should match like any other quantity -- but the observed side is evaluated at the point estimate, so it carries no parameter uncertainty while the generated side does.