




Marginal distribution of each variable. Blue: values generated from the fitted model (200 draws). Black: the observed data. A shift in location or width means the model does not reproduce the marginal, but any question of dynamics arises. LogLik is each row's log-likelihood, and its distribution should match like any other quantity -- but the observed side is evaluated at the point estimate, so it carries no parameter uncertainty while the generated side does.