




Probability integral transform — for each observation, the fraction of the 200 generated values falling below it. Under a correctly calibrated model this is uniform. A hump in the middle means the prediction distribution is too wide, peaks at both ends mean it is too narrow, and a tilt means a location bias. Grey lines: where 95% of bin counts should fall under uniformity. LogLik is each row's log-likelihood and its distribution should match like any other quantity — but the observed side is evaluated at the point estimate, so it carries no parameter uncertainty while the generated side does.