




The same exceedances as a rate, in ten equal-count bins of predicted median, so a misfit concentration among high or low predictions is visible. Dashed line: the nominal 5%. Bars: 95% Wilson interval for each bin. LogLik is each row's log-likelihood, and its distribution should match like any other quantity -- but the observed side is evaluated at the point estimate, so it carries no parameter uncertainty while the generated side does.