




Each observation against the central 95% predictive interval for it, ordered along the x axis by predicted median so the interval reads as a band. Orange points fall outside their own interval; about 5% should. LogLik is each row's log-likelihood, and its distribution should match like any other quantity -- but the observed side is evaluated at the point estimate, so it carries no parameter uncertainty while the generated side does.