




Mean PIT per subject, sorted; 0.5 is correct, and subjects at the ends are the ones the model fits worst. Grey band: the range expected if a subject's own observations were independent. They are not random effect or a persistent state makes them agree -- so the true band is wider, and this one is a guide rather than a test. LogLik is each row's log-likelihood, and its distribution should match like any other quantity -- but the observed side is evaluated at the point estimate, so it carries no parameter uncertainty while the generated side does.